

One Ocean, Many Tasks: Supplementary Appendix

Shuaijun Liu¹, Qifu Wen^{2,3}, Xingwei Chen¹, Xinwen Hu⁴,
Feiyang You¹, Sijia Chen¹, Ningxin Su^{1,*}

¹The Hong Kong University of Science and Technology (Guangzhou)

²Boston University ³Shanghai Jiao Tong University ⁴Tencent

*Corresponding author: ningxinsu@hkust-gz.edu.cn

This standalone appendix expands the data, task, implementation, and evaluation contracts summarized in the main paper. It also reports diagnostic results that are useful for reproducibility and failure analysis but are not needed to establish the main claims.

A Environment Product Contract and Validation

Sources, harmonization, and schema

OceanEnv combines bathymetry from the General Bathymetric Chart of the Oceans (GEBCO) (GEBCO Compilation Group 2025) with horizontal currents from the Copernicus Marine Service Global Ocean Physics Analysis and Forecast (GOPAF) product (Copernicus Marine Service 2024; Copernicus Marine Environment Monitoring Service (CMEMS) 2025). The resulting Network Common Data Form (NetCDF) product is harmonized by cropping the requested region, retaining returned time and depth coordinates, normalizing coordinate names and units, interpolating variables onto one latitude–longitude grid, preserving invalid terrain and NoData cells, and recording the transformations. Harmonization does not mean inventing unavailable depth layers.

The storage merge fills masked elevation cells with zero, so elevation must always be interpreted together with the validity mask. A zero at an invalid cell is not a physical seabed measurement. In the core simulation frame, $(x, y, z) = (\text{east}, \text{depth-down}, \text{north})$, and the dataset vector (u, v) maps to $(u, 0, v)$. Longitude and latitude are converted to local metric x and z about the recorded lower-left geographic origin.

Depth and tide treatment

The realized `tiny` request contains one level at 0.494 m. The `scene` and `public` requests cover 0–200 m and retain 26 source levels from 0.494 to 186.126 m. The pipeline records returned coordinates rather than relabeling a nearest source level as the requested value. Consequently, `tiny` is a time-varying 2D or 2.5D current product, whereas the other two products provide measured horizontal current at multiple discrete depths.

Only the current `scene` release includes tide components. Hourly tide samples are aligned to the base product with nearest-time selection and attenuated as $\exp(-d/z_0)$ with

$z_0 = 50$ m before forming the total current pair. This is a documented engineering approximation chosen for data volume and simulation convenience. It discards high-frequency structure on coarser time grids and is neither a validated 3D tidal model nor suitable for studies that require phase-accurate tidal dynamics.

Validator evidence

The release validator checks required fields, aliases, dimensions, monotonic coordinates, units, finite sampled values, mask semantics, provenance, and optional round-trip fidelity between NetCDF and a NumPy archive (NPZ). Table 2 summarizes the July 21, 2026 audit of the already-generated products.

The legacy warnings identify missing depth or elevation units, total-current units, or mask flag meanings depending on the variant. Regeneration with the updated pipeline writes these attributes. The audited `tiny` slice preserved coordinates and currents at absolute tolerance 10^{-6} .

Portable provenance

Every regenerated NetCDF records source product identifiers, requested and actual region, time, and depth, tide inclusion and depth profile, interpolation and missing-value handling, and generation software. Companion variant manifests use paths relative to the variant directory; private workspace paths are excluded. This contract allows a custom product to fail early before an episode is launched and lets every run retain the selected time and depth indices.

B Canonical Task Contracts

Every task declares an objective, observation contract, feasibility constraints, success condition, and primary metric. Shared observations include pose and relative target cues, local current, and task-specific probes; shared feasibility checks include domain bounds, terrain validity, seafloor clearance, and configured team size.

Go-to-Goal Under Currents (G2G). Reach a goal while rejecting data-grounded drift; success requires entering the goal tolerance, and final goal error is primary.

Station Keeping (SK). Hold position and heading under disturbance; station RMS error measures sustained deviation.




Figure 1: Conceptual view of the environmental channels and ecological signals represented by, or targeted for coupling with, OneOcean. It illustrates terrain, currents, tidal velocity, thermohaline variables, weather, biodiversity, and microplastics; it is not a scale diagram and does not imply that every release variant contains every depicted field.

Canonical field	Dimensions	Unit or Encoding	Contract
Latitude	H	degrees north	Strictly monotonic geographic coordinate.
Longitude	W	degrees east	Strictly monotonic geographic coordinate.
Time	T	Climate and Forecast time unit	Monotonic timestamps retained from the source product.
Depth	D	m, positive down	Actual source levels; a one-level product is valid but not a resolved 3D flow.
Eastward and northward current	$T \times D \times H \times W$	m s^{-1}	Base pair (u, v) or tide-combined pair $(u_{\text{total}}, v_{\text{total}})$.
Elevation	$H \times W$	m relative to sea surface	Bathymetric elevation aligned to the current grid.
Validity mask	$H \times W$	binary	Invalid terrain, crop-exterior, or NoData cells; not necessarily a coastline classifier.

Table 1: Canonical environment-product schema. The validator accepts documented aliases but maps every accepted product to these semantics.

Variant	Current pair	Errors	Warn.	Round-trip
tiny	uo, vo	0	3	Exact
scene	utotal, vttotal	0	4	Not run
public	uo, vo	0	2	Not run

Table 2: Environment-product validation. All variants are PASS_WITH_WARNINGS; warnings concern legacy meta-data rather than numeric or schema errors. Round-trip denotes the fixed tiny ($t = 0, d = 0$) NetCDF-to-NPZ check.

Route Following (RF). Visit an ordered waypoint route; cross-track error and reached waypoints expose partial progress.

Depth-Profile Tracking (DPT). Follow a target depth profile during motion; depth-tracking error is primary.

Area Scan (AS). Cover a bounded region over 3D terrain;

grid coverage fraction measures progress.

Pipeline Inspection (PI). Follow a pipeline corridor and encounter leak cues; the inspection score combines route and leak progress.

Formation Transit (FT). Move a team of 2–10 vehicles while maintaining relative geometry; formation error and collision ratio capture coordination quality.

Surface Pollution Cleanup (SPC). Assign 2–10 vehicles to discrete surface sources and satisfy per-source dwell requirements; cleanup rate is completed-source fraction.

Underwater Pollution Lift (UPL). Coordinate five vehicles around a payload and lift it to the surface; lift completion defines success.

Fish Protection and Herding (FH). Use eight vehicles to contain and escort a fish school under currents; containment and escort completion define success.

Exact progress metrics

Several task metrics are normalized ratios with distinct denominators. Let $\mathcal{V}$ be the visited scan cells in an $H \times W$ grid, N_{det} the detected leaks among N_{leak} placed leaks, and N_{done} the sources whose dwell requirement is complete among N_{src} sources. The released evaluator computes

$$\text{Cov}_{\text{AS}} = \frac{|\mathcal{V}|}{HW}, \quad (1)$$

$$\text{Score}_{\text{PI}} = \frac{N_{\text{det}}}{N_{\text{leak}}}, \quad (2)$$

$$\text{Rate}_{\text{SPC}} = \frac{N_{\text{done}}}{N_{\text{src}}}, \quad (3)$$

$$e_{\text{FT}} = \frac{1}{N} \sum_{i=1}^N \|p_i - (p_g + r_i)\|_2, \quad (4)$$

$$\text{Prog}_{\text{FH}}(t) = \text{clip}\left(1 - \frac{d_t}{d_0}, 0, 1\right). \quad (5)$$

Here $p_g + r_i$ is agent i 's assigned formation goal, while d_t and d_0 are the current and initial fish-centroid distances to the escort goal. These quantities expose partial progress but are not interchangeable physical measurements.

Tiers 1, 2, and 3 preserve these contracts while jointly changing current gain, tolerances, route or coverage extent, sensing difficulty, and coordination tightness. They are therefore named presets rather than a claim that one scalar fully orders task difficulty. Navigation and tracking tasks primarily stress disturbance rejection, sensing and inspection tasks stress spatial coverage and event encounters, and coordination and pollution tasks add team geometry, assignment, or dwell-time coupling.

C Implementation Defaults, Pollution, and ROS 2

Serialized engineering defaults

The main paper defines current coupling, action clipping, diagonal engineering dynamics, and terrain feasibility. Table 3 records the numerical defaults for the six-degree-of-freedom (6DoF) mode serialized in the released run specifications so that those equations can be reproduced without relying on hidden simulator state.

Each run also records whether it uses direct kinematic integration, the diagonal 3DoF mode, or the bounded-attitude 6DoF mode, preventing mode selection from being inferred from output-directory names.

Pollution fields and cleanup semantics

The analytic probe field is a current-advected Gaussian,

$$C(p, t) = m(t) C_{\text{peak}} \exp[-\|p - \mu(t)\|^2 / (2\sigma^2)], \quad (6)$$

with default $\sigma = 35$ m and $C_{\text{peak}} = 1$. An agent in the 35 m continuous-containment neighborhood applies $m(t + \Delta t) = m(t) \max(0, 1 - 0.12\Delta t)$.

Three-dimensional transport and transformation model

The Ocean Current Prediction Network (OCPNet) pollution engine evolves each configured species s through the implemented balance

$$\frac{\partial C_s}{\partial t} = -\mathbf{u} \cdot \nabla C_s + \nabla \cdot (K_s \nabla C_s) - \lambda_s C_s + \sum_r \nu_{sr} R_r + \mathcal{U}_s + \mathcal{P}_s + S_s - Q_s, \quad (7)$$

$$R_r = k_{r0} \exp[\alpha_r (T - 298.15)] f_r(\text{pH}) \times \prod_{j \in \mathcal{R}_r} C_j, \quad (8)$$

$$f_r(\text{pH}) = \max(1 + a_r \Delta \text{pH} + b_r \Delta \text{pH}^2, 0).$$

Here $\Delta \text{pH} = \text{pH} - 7$, C_s is concentration, $\mathbf{u} = (u, v, w)$ is velocity, K_s is diffusivity, λ_s is first-order decay, and ν_{sr} is a signed stoichiometric coefficient. The optional terms $\mathcal{U}_s$ and $\mathcal{P}_s$ represent biological uptake and precipitation, while S_s and Q_s collect externally configured sources and sinks.

The code computes spatially varying diffusivity from temperature T , wave velocity U_w , and salinity S as

$$\begin{aligned} K_T &= K_{s0} [1 + 0.02(T - 273.15)], \\ K_w &= K_T + 0.1U_w, \\ K_s &= K_w [1 + 0.05(1 - S/35)]. \end{aligned} \quad (9)$$

These coefficients reproduce the released numerical parameterization rather than a calibrated turbulence closure. When decay is enabled, its implemented environmental dependence is

$$\lambda_s = \lambda_{s0} \exp\left(-\frac{E_{a,s}}{RT}\right) f_s(\text{pH}) f_s(O), \quad (10)$$

$$f_s(\text{pH}) = \max(1 + a_s \Delta \text{pH} + b_s \Delta \text{pH}^2, 0),$$

$$f_s(O) = \max(1 + c_s O + d_s O^2, 0),$$

where the exponential factor is omitted when $E_{a,s} = 0$, R is the gas constant, and O is dissolved oxygen. Biological uptake, when configured, contributes the Michaelis–Menten sink

$$\mathcal{U}_s = -V_{\text{max},s} e^{\alpha_{T,s}(T-298.15)} (1 + \alpha_{L,s} L) \frac{C_s}{K_{1/2,s} + C_s} B, \quad (11)$$

with light L and phytoplankton biomass B . For a configured cation–anion precipitation pair, the saturation and rate terms are

$$\begin{aligned} \text{SI} &= \frac{C_M C_A}{K_{\text{sp}}}, \\ r_p &= k_p e^{\alpha_p (T-298.15)} f_p(\text{pH}) \max(\text{SI} - 1, 0), \end{aligned} \quad (12)$$

which subtract r_p from both dissolved species and add it to the configured precipitate before its vertical settling term is applied.

For transport, the implementation uses a first-order upwind difference. Along x , for example,

$$(\delta_x^{\text{up}} C)_i = \begin{cases} (C_i - C_{i-1})/\Delta x, & u_i \geq 0, \\ (C_{i+1} - C_i)/\Delta x, & u_i < 0, \end{cases} \quad (13)$$

Field	Default	Meaning	Field	Default	Meaning
Δt	1.0 s	Outer step	$v_{\max}$	1.2 m s^{-1}	Relative speed cap
M_ℓ	(12, 12, 12)	Linear inertia proxy	M_a	(6, 6, 6)	Angular inertia proxy
D_ℓ	(8, 8, 8)	Linear damping	D_a	(3, 3, 3)	Angular damping
K_ℓ	(18, 18, 18)	Linear tracking gain	K_a	(10, 10, 10)	Angular tracking gain
K_d	(2, 2, 2)	Angular damping gain	$\omega_{\max}$	(1.2, 1.2, 1.2) rad s^{-1}	Angular-rate cap
k_{yaw}	2.0	Yaw-rate gain	k_{att}	2.0	Roll and pitch stabilization
$\theta_{\max}$	0.6 rad	Roll and pitch limit			

Table 3: Default 6DoF engineering parameters serialized in each run specification.

and the centered flux form of the corresponding diffusion term is

$$(\mathcal{D}_x C)_i = \frac{K_i(C_{i+1} - C_i) - K_{i-1}(C_i - C_{i-1})}{\Delta x^2}. \quad (14)$$

The same constructions are applied along y and z , and explicit Euler combines every active process:

$$\begin{aligned} C_s^{n+1} &= C_s^n + \Delta t \mathcal{F}_s^n, \\ \mathcal{F}_s &= \mathcal{A}_s + \mathcal{D}_s - \lambda_s C_s + \sum_r \nu_{sr} R_r \\ &\quad + \mathcal{U}_s + \mathcal{P}_s + S_s - Q_s. \end{aligned} \quad (15)$$

The solver rejects velocity fields that violate

$$\max\left(\frac{\|u\|_\infty \Delta t}{\Delta x}, \frac{\|v\|_\infty \Delta t}{\Delta y}, \frac{\|w\|_\infty \Delta t}{\Delta z}\right) \leq 1, \quad (16)$$

and the diffusion module exposes the conservative helper bound $\Delta t_{\text{diff}} = \frac{1}{2} \min(\Delta x, \Delta y, \Delta z)^2 / \|K_s\|_\infty$. Dirichlet, Neumann, periodic, and open boundaries are configurable. The general three-dimensional update does not clip negative concentrations, so stable timesteps and physically consistent process parameters remain the user’s responsibility.

The benchmark wrapper uses a $28 \times 28 \times 10$ grid, a 2 s internal step, base diffusivity 8×10^{-8} , point-source emission rate 0.02, 8 m sink radius, and sink strength 0.15 s^{-1} . Its default benchmark case activates transport, diffusion, the point source, and the robot sink; the other transformation terms above are modular capabilities and are not silently activated. Dataset horizontal current is resampled and repeated over model depth with zero vertical current. This is an explicit approximation, not a measured three-dimensional velocity profile. The robot sink multiplies cells in its local grid neighborhood by $\max(0, 1 - 0.15\Delta t)$ at each internal update.

Dataset-driven surface multispecies proxy The supplementary visualization uses a separate two-dimensional proxy driven by a selected depth slice of the dataset currents. For any field F , one transport substep is

$$\mathcal{T}_n(F) = F + \Delta t [-u_n F_x - v_n F_y + \kappa(F_{xx} + F_{yy})]. \quad (17)$$

Let A , B , and C denote normalized microplastic, crude-oil, and aggregate fields, respectively. Transport first gives $A^* = \mathcal{T}_n(A^n)$, $B^* = \mathcal{T}_n(B^n)$, and $C^* = \mathcal{T}_n(C^n)$. With $[x]_+ = \max(x, 0)$, aggregation and weathering are then

$$r_a = k_a[A^*]_+[B^*]_+, \quad r_w = k_w[B^*]_+, \quad (18)$$

Interface	ROS type and meaning
/agent_NNN/odom	nav_msgs/Odometry; ENU pose and velocity
/agent_NNN/current	Vector3Stamped; ENU current
/agent_NNN/pollution	Float32; local scalar probe
/agent_NNN/bathymetry	Float32; positive water depth
/agent_NNN/cmd_vel	subscribed ENU velocity command
/clock, /episode/metrics	simulation time and JavaScript Object Notation episode state
/reset, /step, /get_task_state	std_srvs/Trigger services

Table 4: ROS 2 interface. Topic prefixes are /oneocean/agents; agent IDs are zero-padded.

and the implemented update before masking and normalization is

$$\begin{aligned} \tilde{A} &= [A^* - \Delta t r_a]_+ + 0.010 S_A, \\ \tilde{B} &= [B^* - \Delta t(r_a + r_w)]_+ + 0.008 S_B, \\ \tilde{C} &= [C^* + \Delta t r_a]_+. \end{aligned} \quad (19)$$

Gaussian source maps S_A and S_B continuously replenish the two parent fields. Land-mask cells are then set to zero, and every nonzero species field is independently divided by its spatial maximum. Consequently, Figure 2 shows relative plume shape and reaction overlap rather than calibrated concentration or mass-conserving forecasts. This surface proxy uses centered numerical gradients and must not be conflated with the upwind three-dimensional solver in Equations 7–16.

Building on the distinct localization, containment, and discrete-cleanup semantics defined in the main paper, the exact canonical SPC presets are: 12 sources, 8 s dwell, and one required agent at medium difficulty; 6 sources, 6 s dwell, and two required agents at hard difficulty.

Optional Robot Operating System 2 (ROS 2) Jazzy bridge

The bridge maps core $(x, y, z) = (\text{east}, \text{down}, \text{north})$ to the ROS east-north-up (ENU) frame $(x, z, -y)$. External commands traverse the same clipping, dynamics, current, constraints, recorder, and task-update path as built-in controllers; internal-control mode can publish observations without accepting external motion commands.

The package was built and launched in a ROS 2 Jazzy container; topic discovery, one externally commanded step, state




Figure 2: Six sampled times from the dataset-driven surface multispecies proxy. The overlays show normalized microplastic, crude-oil, and aggregate fields together with aggregation and oil-weathering rates. The land mask blocks the proxy fields. Because each species is normalized independently after every substep, this panel is qualitative evidence of transport and reaction structure rather than a mass forecast.

query, and reset completed successfully. This is an interface smoke test, not a ROS performance result, and ROS remains an optional dependency outside the core requirements.

D Metrics and Reproducibility Protocol

Metric definitions

The main paper defines success rate (SR), timeout-aware episode time T , and the action-energy proxy E . The additional team-level aggregate used in the result tables is the near-collision ratio

$$\text{Coll} = \frac{1}{T_{\text{ep}}} \sum_{t=1}^{T_{\text{ep}}} \mathbf{1} \left[\min_{i < j} \|p_t^{(i)} - p_t^{(j)}\|_2 \leq r_{\text{coll}} \right]. \quad (20)$$

Task-primary metrics follow the per-task contracts in Appendix Sec. B and are not pooled across incompatible physical meanings.

Supplementary diagnostic budgets

All diagnostics reuse the same serialized tier presets, controller settings, dataset indices, and environment provenance as their corresponding headline study; only the diagnostic factors named in Table 5 are varied.

Compute and run manifests

The main paper reports the compute platform. For provenance, each run stores an integer seed, normalized configuration, task and dynamics specification snapshots, environment-product indices, software and Git state, episode metrics, and recorder streams. Aggregation scripts consume these artifacts to regenerate comma-separated-value and Markdown tables rather than transcribing console output.

The reproducibility boundary is explicit: data validation and ROS adapter tests are release tests; they are not added to performance averages. Likewise, the three-seed planner audit and three-episode-per-scene pollution comparison diagnose mechanisms but do not replace the larger headline budgets.

Model boundaries and extension priorities

The pollution fields are controlled sensing signals rather than in-situ labels or validated spill forecasts. Surface Pollution Cleanup completes discrete sources after a dwell requirement, whereas Underwater Pollution Lift removes mass from a concentration field; their cleanup quantities therefore have different semantics. Tide attenuation remains an engineering approximation rather than a phase-accurate three-dimensional tide model. The ROS 2 bridge has passed message, reset, frame-conversion, and command smoke tests,

Diagnostic	Budget	Configuration	Purpose
Planner audit	3 seeds per task and method	72 episodes; stride 200; 96-token cap	Validity, fallback, completion, and failure diagnosis.
Dynamics diagnostic	20 episodes per task and difficulty	10 tasks; medium and hard; three dynamics modes	Sensitivity to engineering dynamics.
Current-gain sweep	5 seeds per task and method	300 episodes; hard; 10 tasks; $g \in \{0, 1, 2\}$	Data-current disturbance sensitivity.
Synthetic-tide stress	10 seeds per task and method	200 episodes; hard; 10 tasks; 0.45 m s^{-1} amplitude	Periodic-disturbance sensitivity.
External pollution diagnostic	3 episodes per scene and configuration	Four scenes; two currents; two field models	Localization-model comparison.
External scaling	10 episodes per scene and team size	80 episodes; two scenes; $N \in \{2, 4, 8, 10\}$	Scene-backed team-size sensitivity.

Table 5: Budgets for supplementary diagnostic studies. Headline experiment budgets remain attached to their result tables in the main paper.

Hard task	Kinematic SR	3DoF SR	6DoF SR
Formation Transit	55%	50%	50%
Surface Cleanup	95%	85%	85%
Pollution Lift	100%	75%	75%
All-task macro	85.5%	81.5%	81.5%

Table 6: Selected dynamics sensitivity at hard difficulty (20 episodes per task and mode).

but it has not been benchmarked for latency, determinism, middleware load, or hardware-in-the-loop performance. The highest-priority extensions are measured three-dimensional flow, identified vehicle coupling through computational fluid dynamics (CFD) or system identification, physically validated pollutant transport, and ROS-in-the-loop experiments. These additions can preserve the current task and reporting contracts while increasing physical realism.

E Additional Core-Backend Diagnostics

Dynamics-mode sensitivity

Table 6 selects the hard-setting tasks with the largest observed dynamics sensitivity from a 20-episode-per-task diagnostic.

Across all ten hard tasks, macro SR is 85.5% for kinematic and 81.5% for both 3DoF and 6DoF. The identical 3DoF and 6DoF SRs under this controller do not establish physical equivalence; they show that bounded roll and pitch did not change terminal success in this diagnostic, whereas replacing the velocity state with direct kinematic integration did.

Current and synthetic-tide stress

Heuristic SR decreases from 48.0% without data-current forcing to 8.0% at doubled current, while its energy and collision exposure increase. BC is not monotonic at the smaller gains, but reaches its lowest SR at $g = 2$. The synthetic-tide stress also separates the methods, with 17.0% SR for the heuristic and 1.0% for BC. Because the current-gain cells contain only five episodes per task, these results establish a sensitivity trend rather than precise task-level rankings.

Planner validity, fallback, and cost

The instrumented audit contains three medium-difficulty seeds for each of eight methods and three tasks, for 72 episodes total. Tables 8 and 9 aggregate each large language

Stress	Method	N	SR	E	Coll.
Data current, $g = 0$	Heuristic	50	48.0	3232.5	5.80
Data current, $g = 1$	Heuristic	50	20.0	3371.6	8.24
Data current, $g = 2$	Heuristic	50	8.0	3775.0	9.86
Data current, $g = 0$	BC	50	6.0	1546.8	4.14
Data current, $g = 1$	BC	50	12.0	1807.1	4.96
Data current, $g = 2$	BC	50	4.0	2323.3	2.44
Synthetic tide	Heuristic	100	17.0	3489.2	6.29
Synthetic tide	BC	100	1.0	1557.6	6.49

Table 7: Archived hard-setting disturbance diagnostics across all ten tasks. SR and Coll. are percentages. Each data-current row contains five episodes per task. The synthetic-tide rows contain ten episodes per task, use $g = 0$, and apply a 0.45 m s^{-1} sinusoid with a 300s period. These focused sweeps diagnose sensitivity and are separate from the headline budgets.

model (LLM) across its nine episodes. Valid-plan ratio is weighted by planner calls; latency and tokens are weighted by uncached calls.

LLM	SR	Calls	Valid	Fallback	Comp.
ChatGLM3-6B	0.556	49	0.367	0.633	0.758
Llama2-7B	0.333	60	0.950	0.050	0.736
Llama3-8B	0.444	61	0.869	0.131	0.786
Mistral-7B	0.556	55	0.800	0.200	0.804
Qwen2-7B	0.556	43	0.814	0.186	0.790
Qwen2.5-7B	0.667	59	0.864	0.136	0.809

Table 8: Per-model planning outcomes across SPC, AS, and PI (nine episodes per LLM).

Model-level and task-level rankings also diverge: Llama2-7B attains 0.950 validity but only 0.333 SR, whereas ChatGLM3-6B attains 0.367 validity and 0.556 SR. At task level, cleanup’s low valid ratio coexists with high completion because invalid or duplicate assignments can fall back to a competent deterministic allocator; AS can revisit covered cells despite valid syntax, and PI still requires physical leak encounters between low-frequency plans. Planner validity is therefore a mechanism diagnostic, not a substitute for task success.

LLM	Latency ms per call	Prompt tokens per call	Output tokens per call
ChatGLM3-6B	1466.1	628.1	76.7
Llama2-7B	7429.6	633.7	728.7
Llama3-8B	2068.1	454.1	475.9
Mistral-7B	5150.6	616.5	676.7
Qwen2-7B	2370.6	603.3	628.6
Qwen2.5-7B	2184.5	600.6	623.9

Table 9: Uncached per-call planning cost in the same audit.

Task	SR	Valid	Fallback	Latency ms per call	Completion
SPC	0.833	0.430	0.570	3199.0	0.972
AS	0.278	0.831	0.169	3602.0	0.758
PI	0.444	0.891	0.109	3337.1	0.611

Table 10: Cross-model planner diagnostics by task (six LLMs, three seeds each).

F Additional External-Backend Diagnostics

The external pollution study compares an analytic advected Gaussian probe with OCPNet-3D in Dam, OpenWater, PierHarbor, and SimpleUnderwater. Each combination of scene, model, and current setting uses three localization episodes; Table 11 therefore aggregates 12 episodes per row. It is a focused observation-model diagnostic rather than a high-confidence performance ranking.

OCPNet-3D lowers mean localization error by 5.12 m at nominal current and 3.36 m under stronger current, while matching or improving aggregate success. This result supports the narrower claim that richer regeneratable concentration structure can change sensing difficulty; it does not validate OCPNet as a real spill forecast. Containment outputs are not collapsed into the same table because analytic remaining-mass state and OCPNet grid mass have different scales and should not be compared as a shared physical quantity.

Scene-backed cleanup scaling

Increasing the team from two to four agents raises aggregate success from 40% to 100% and approximately doubles removed mass. The smaller gains from four to ten agents indicate saturation under the fixed 15 s containment horizon. Removed mass is comparable within this analytic-field sweep only and is not equated with OCPNet grid mass.

The route-transfer results retained in the main paper reveal substantial scene dependence despite a shared task contract. Together, the route and pollution diagnostics motivate reporting scene, current scale, dataset clock, controller profile, and observation model in every external-backend manifest rather than treating “external simulator” as a single condition.

Scale	Probe	N	SR (%)	Error (m)	Δ error (m)
1.0	Analytic	12	41.7	25.93	0.00
1.0	OCPNet-3D	12	50.0	20.81	−5.12
2.5	Analytic	12	50.0	26.00	0.00
2.5	OCPNet-3D	12	50.0	22.64	−3.36

Table 11: Pollution-localization diagnostic over four scenes (12 episodes per row). Δ error is measured against the analytic probe at the same current scale.

Agents	Episodes	SR (%)	Removed mass (kg)
2	20	40.0	0.2081
4	20	100.0	0.4302
8	20	100.0	0.4817
10	20	100.0	0.4886

Table 12: Analytic-containment scaling at nominal current, aggregated over OpenWater and PierHarbor with ten episodes per scene and team size.

References

- Copernicus Marine Environment Monitoring Service (CMEMS). 2025. Global Ocean Physics Analysis and Forecast. <https://data.marine.copernicus.eu/>. Accessed: 2025-09-07.
- Copernicus Marine Service. 2024. Copernicus Marine Service: Ocean Monitoring and Forecasting.
- GEBCO Compilation Group. 2025. The GEBCO 2025 Grid – A Global 15 Arc-Second Bathymetry Dataset.