

One Ocean, Many Tasks: A Data-Grounded Simulation Suite and Benchmark for Marine Robotics

Shuaijun Liu¹, Qifu Wen^{2,3}, Xingwei Chen¹, Xinwen Hu⁴, Feiyang You¹, Sijia Chen¹, Ningxin Su^{1,*}

¹The Hong Kong University of Science and Technology (Guangzhou) ²Boston University

³Shanghai Jiao Tong University ⁴Tencent *Corresponding author: ningxinsu@hkust-gz.edu.cn




Figure 1: OneOcean overview. A shared ocean dataset supplies two simulation backends and a ten-task benchmark covering navigation, sensing, coordination, and pollution response.

Abstract

Simulation is essential for marine robotics, yet existing platforms often trade off visual realism, experimental throughput, and ocean conditions derived from real data. We present OneOcean, a data-grounded simulation suite and benchmark that combines bathymetry, currents, and tides in a shared spatiotemporal dataset. Optional pollution fields support pollution-response evaluation. OneOcean defines ten tasks covering navigation, sensing, coordination, and pollution response, with teams of 2–10 vehicles. The simulator supports three to six Degrees of Freedom (DoF) and records standardized metrics and run configurations across tasks, difficulty tiers, and team sizes. Experiments across dataset variants and underwater scenes compare heuristic control, Behavior Cloning (BC), and Large Language Model (LLM)-assisted planning under ocean currents.

[Project](#) [Code](#) [Demo](#) [Hugging Face](#)

Introduction

Marine robots must operate under complex, highly variable ocean conditions. In practice, **currents and tides** can dominate vehicle motion, while field trials remain expensive and time-limited and at-sea sensing and communication are constrained (Zereik et al. 2018). Simulation is therefore essential,

but existing platforms often trade off between (i) visually realistic underwater worlds, (ii) high-throughput, reproducible experimentation, and (iii) **ocean-condition realism grounded in real data**.

We present **OneOcean**, a data-grounded simulation suite and benchmark for marine robotics. OneOcean is built around a unified spatiotemporal ocean-environment dataset in Network Common Data Form (NetCDF). It harmonizes bathymetry from the General Bathymetric Chart of the Oceans (GEBCO) (GEBCO Compilation Group 2025) with data-driven currents and tides from the Copernicus Marine Service Global Ocean Physics Analysis and Forecast (GOPAF) product (Copernicus Marine Service 2024; Copernicus Marine Environment Monitoring Service (CMEMS) 2025). OneOcean defines a **10-task benchmark** spanning single-agent navigation and tracking, spatial sensing and coverage, and multi-agent coordination and pollution-related tasks, with support for teams of 2–10 vehicles. The dataset variants support local iteration, scene construction, and public release under common task definitions. Figure 1 summarizes this organization.

OneOcean provides two complementary execution backends that share task definitions and evaluation semantics. The *core benchmark backend* runs deterministic episodes at scale and records standardized metrics and run configurations. The *external underwater-scene backend* runs the same tasks in

Symbols: ● Provided; ● Not provided. Tasks: G2G, SK, RF, DPT, AS, PI, FT, SPC, UPL, and FH.

Platform or Work	Venue	Real-data In Loop	Multi- agent	Cross- backend	3D Currents	Pollution or Spill	Tide Modeling	Canonical Task	Metrics Protocol	Headless Batch	G2G	SK	RF	DPT	AS	PI	FT	SPC	UPL	FH
MIMIR-UW	2023	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●
HoloOcean 2.0	2025	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●
OceanSim	2025	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●
MarineGym	IROS 2025	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●
OceanGym	2025	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●
URLab	ICRA 2026	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●
DREAM	ICRA 2026	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●
Douro MARL	ICRA 2026	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●
Map-TIDAL	ICRA 2026	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●
CAMAR	AAAI 2026	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●
RobotArena ∞	ICLR 2026	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●
CaP-X	ICML 2026	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●
OneOcean (ours)	2026	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●	●

Table 1: System-level comparison and task coverage. Entries are compiled from the cited papers and public documentation (Álvarez-Tuñón et al. 2023; Romrell et al. 2025; Song et al. 2025; Chu et al. 2025; Xue et al. 2025; Embley-Riches et al. 2026; Wu et al. 2026; Dal Fabbro et al. 2026; Ruediger and Banerjee 2026; Pshenitsyn, Panov, and Skrynnik 2026; Jangir et al. 2026; Fu et al. 2026).

packaged underwater scenes from an existing simulator using OneOcean current exports. In summary, our contributions are:

- A shared ocean-environment dataset and tools for exporting currents to simulators, with variants for fast iteration, scene construction, and public release.
- A 10-task benchmark with difficulty tiers, standardized metrics, reproducible run records, and 3–6 DoF vehicle dynamics, with 6DoF used for the main quantitative evaluation.
- Experiments across the two backends that reveal task difficulty, scaling, and planner-compute trade-offs under data-grounded disturbances.

Related Work

Classical simulators built on Gazebo (Koenig and Howard 2004), including UUV Simulator (Manhães et al. 2016) and UWSim (Prats et al. 2012), provide dynamics and sensor stacks for robotics workflows. More recent systems target higher-fidelity rendering or higher-throughput experimentation, including MARUS (Lončar et al. 2022), OceanSim on Isaac Sim (Song et al. 2025; Makovychuk et al. 2021), DAVE (Zhang et al. 2022), MarineGym (Chu et al. 2025), and URLab, which integrates Unreal Engine with MuJoCo (Embley-Riches et al. 2026). OneOcean supplies a shared ocean dataset, a ten-task benchmark, and a reporting protocol across two execution backends. Table 1 compares system capabilities and task coverage.

Ocean-condition realism is commonly provided by numerical models and operational forecasting systems. Representative community models include Delft3D (Lesser et al. 2004), NEMO (Madec and the NEMO Team 2019), MITgcm (Strobach et al. 2018), ROMS (Shchepetkin and McWilliams 2005), and HYCOM (HYCOM Consortium for Data-Assimilative Ocean Modeling 2024). Operational products from the Copernicus Marine Service, including GOPAF, provide data-assimilative currents and hydrographic variables. GEBCO provides global gridded seafloor elevation, while

NetCDF stores labeled multidimensional coordinates, variables, units, and provenance in a self-describing representation. These sources are public, globally scoped, and machine-queryable, allowing the same ocean condition to be extracted and reproduced across runs. Toolkits such as Parcels (Delandmeter and van Sebille 2019) and OpenDrift (Dagestad, Röhrs et al. 2018) support Lagrangian drift and dispersion analysis. OneOcean uses these ocean data sources to define reproducible robotics tasks across simulation backends.

Recent marine-robotics studies address monitoring, inspection, and plume mapping. DREAM studies vision-language monitoring through simulation and tank trials (Wu et al. 2026); Map-TIDAL allocates cooperative underwater-inspection tasks (Ruediger and Banerjee 2026); and Douro MARL maps a data-driven river plume with multiple vehicles (Dal Fabbro et al. 2026). Recent robotics benchmarks evaluate multi-agent routing and transferable control. CAMAR defines high-throughput continuous multi-agent routing and a three-tier protocol (Pshenitsyn, Panov, and Skrynnik 2026); RobotArena ∞ converts real demonstrations into scalable simulation tests for robot policies (Jangir et al. 2026); and CaP-X evaluates executable embodied agents across multiple manipulation suites (Fu et al. 2026). OneOcean combines marine task definitions with ocean data, reproducible run records, and a common evaluation protocol across two backends.

OneOcean Suite

OneOcean combines a shared ocean dataset, two simulation backends, and a ten-task evaluation protocol.

Unified spatiotemporal ocean dataset

The shared NetCDF representation harmonizes GEBCO bathymetry with GOPAF currents and optional tides. It exposes named latitude, longitude, time, and depth coordinates, eastward and northward current components, elevation, and a mask for invalid terrain and NoData cells. The pipeline crops the requested region, retains the returned time and depth levels, aligns variables on one horizontal grid, and records provenance.

The core backend selects a fixed (t, d) slice for deterministic

Variant	Region and period	$T \times D \times H \times W$	Resolution; size; use
tiny	42.1–42.7 N, 71.2–70.2 W; 2025	$357 \times 1 \times 62 \times 101$	$\approx 0.01^\circ$; 29 MB; tests and web
scene	32–33 N, 66.5–65.5 W; Dec. 2025	$23 \times 26 \times 240 \times 240$	0.00418° ; 690 MB; scenes
public	30–40 N, 72–62 W; 2025	$357 \times 26 \times 41 \times 41$	0.25° ; 175 MB; release

Table 2: Realized combined NetCDF variants. Sizes are the merged products themselves; source subsets and derived assets are distributed separately.




Figure 2: Base current field visualization over bathymetry. The horizontal current components are u_{data} and v_{data} ; surface color indicates speed magnitude (m s^{-1}). View configuration: pitch 25° , azimuth 45° , and depth index 0 (source depth 0.494 m). The terrain surface is colored by speed and overlaid with arrows indicating current direction.

rollouts, whereas the external backend advances a dataset clock. Table 2 reports the realized products: the 0–200 m requests returned 26 levels through 186.13 m, while **tiny** contains one near-surface level at 0.494 m. Appendix Sec. A specifies the full schema, mask semantics, coordinate conventions, validation results, and tide sampling and depth treatment.

The validity mask marks unusable cells. The tide-enabled scene variant aligns tide samples to the dataset clock and applies depth attenuation.

Data-grounded disturbance signals

OneOcean exports either a fixed gridded slice for high-throughput rollouts or a temporal profile for scene-backed execution. Both apply current through a recorded relative-velocity convention in the core frame $(x, y, z) = (\text{east, depth-down, north})$, and every artifact records its extraction indices. Figure 2 shows one current slice over bathymetry; Appendix Sec. A details sampling, export, and provenance.

Environment queries and probes. Shared queries expose local current, bathymetry, mask validity, and task signals. Pollution tasks observe $o_i^c(t) = C(p_i(t), t)$ from either a current-advected Gaussian or an Ocean Current Prediction Network (OCPNet) field. Both models generate reproducible synthetic concentration signals. Appendix Sec. C gives the probe, sink, and discrete-source cleanup semantics.

Category	Task	Agents	Goal	Primary metric
Navigation	G2G	1	reach	final goal error
Navigation	SK	1	hold	station root-mean-square error
Navigation	RF	1	follow	cross-track error
Navigation	DPT	1	track	depth-tracking error
Sensing and Coverage	AS	1–10	scan	coverage fraction
Sensing and Inspection	PI	1–10	inspect	inspection score
Coordination	FT	2–10	transit	formation error
Pollution	SPC	2–10	clean	cleanup rate
Coordination	UPL	5	lift	lift success
Coordination	FH	8	herd	containment success

Table 3: The OneOcean task suite (10 tasks). Each task is evaluated at Tiers 1–3 and paired with a compact goal type and a primary task-specific metric.

Task suite and difficulty tiers

OneOcean defines ten tasks spanning navigation and tracking (G2G, SK, RF, DPT), sensing and inspection (AS, PI), and coordination and pollution response (FT, SPC, UPL, FH). Table 3 gives team sizes and primary metrics; Appendix Sec. B provides the complete objectives, observations, constraints, and success conditions.

Tiers 1, 2, and 3 jointly increase current strength, precision, route and coverage extent, sensing difficulty, and coordination tightness while preserving each task’s interface. Every task declares its objective, observations, feasibility constraints, success condition, and primary metric. The release also includes a lightweight Robot Operating System 2 (ROS 2) bridge for state, probes, and velocity commands; its frames, topics, services, and interface checks appear in Appendix Sec. C.

Navigation and tracking tasks test disturbance rejection and sustained trajectory accuracy; sensing and inspection tasks require spatial coverage or physical cue encounters; coordination and pollution tasks add team geometry, assignment, and dwell-time coupling. We report success together with task-specific progress.

Two execution backends

The two backends share task definitions and reporting formats. The headless **core benchmark backend** consumes a NumPy Archive (NPZ) current slice and task and action configurations. It returns observations, metrics, run records, and recorder streams, and supports 3–6 DoF vehicle dynamics (6DoF in the main experiments) (Manhães et al. 2016). The **external underwater-scene backend** runs the task suite in packaged underwater scenes using a temporal OneOcean current export. It returns route and task metrics alongside simulation recordings.

Benchmark protocol and metrics. Actions specify desired relative velocity, clipped as $a_t = \text{clip}(\tilde{a}_t; v_{\text{max}})$. With current gain g ,

$$c_t = g [u_{\text{data}}(x_t, z_t), 0, v_{\text{data}}(x_t, z_t)] + c_{\text{tide}}(t), \quad (1)$$

where u_{data} and v_{data} are sampled horizontal currents and c_{tide} is an optional synthetic robustness signal. With body velocity $\nu = [\nu_\ell^\top, \nu_a^\top]^\top$, rotation R , and diagonal M, D, K_ℓ, K_a, K_d ,

Method	Macro (Tier 2)			Macro (Tier 3)			Tier 3 key-task metrics					
	Suite SR(%)	E	Coll. (%)	Suite SR(%)	E	Coll. (%)	G2G SR(%)	RF SR(%)	SK RMS	Scan cov.	Pipe score	Cleanup rate
Heuristic	84.6	1842.3	5.68	29.0	3352.4	8.19	74.0	54.0	1.07	0.89	0.78	0.56
BC	58.6	1271.2	9.72	6.0	1635.1	6.41	34.0	0.0	26.52	0.89	0.67	0.06

Table 4: Headline core-benchmark results (6DoF): macro averages at Tiers 2 and 3 plus Tier 3 key-task metrics, using 50 episodes per task and tier. Success Rate (SR) is reported in percent; Macro Suite SR is an unweighted mean across the 10 tasks. Coll is the mean collision-step ratio (%) with a 2.0 m radius. Best values per column are highlighted.

Method	Avg. SR(%)	Cleanup SR(%)	Cleanup rate	Scan SR(%)	Scan cov.	Pipe SR(%)	Pipe score	$T(s)$	E
Heuristic	65.8	60.0	0.87	77.5	0.80	60.0	0.69	776.8	9132.4
BC	45.0	10.0	0.32	67.5	0.79	57.5	0.61	921.6	6908.2
LLM-ChatGLM3-6B	55.0	60.0	0.87	77.5	0.82	27.5	0.54	860.3	9671.5
LLM-Qwen2.5-7B	47.5	57.5	0.81	57.5	0.78	27.5	0.57	1077.2	11703.8
LLM-Qwen2-7B	45.0	57.5	0.86	67.5	0.81	10.0	0.51	991.4	10422.1
LLM-Mistral-7B	42.5	55.0	0.82	55.0	0.78	17.5	0.52	1044.2	11288.9
LLM-Llama3-8B	44.2	67.5	0.84	47.5	0.77	17.5	0.59	1099.1	11614.2
LLM-Llama2-7B	30.8	55.0	0.87	27.5	0.75	10.0	0.46	1156.4	11654.3

Table 5: Core-benchmark planning-suite performance at Tier 2 with doubled currents: heuristic vs BC vs local LLM planners on three planning-sensitive tasks, using 40 episodes per task (120 per method). Best values per column are highlighted.

the diagonal 6DoF update is

$$\begin{aligned}
M\dot{\nu} &= \tau - D\nu, & \tau_\ell &= K_\ell(\nu_\ell^{\text{cmd}} - \nu_\ell), \\
\tau_a &= K_a(\nu_a^{\text{cmd}} - \nu_a) - K_d\nu_a, & \dot{p} &= R\nu_\ell + c_t.
\end{aligned} \quad (2)$$

Commands are transformed to the body frame; yaw follows horizontal motion while roll and pitch are stabilized. Kinematic mode directly integrates $a_t + c_t$, and 3DoF suppresses roll and pitch. The model uses diagonal inertia and damping, bounded attitude, and externally sampled current transport. Added mass, Coriolis and restoring terms, thruster allocation, and two-way fluid coupling are omitted. Hard terrain feasibility is

$$\mathcal{C} = \{(x, y, z) \mid m_{\text{land}}(x, z) < \tau_{\text{land}}, y + \delta_{\text{clr}} \leq -h(x, z)\}, \quad (3)$$

where m_{land} is mask validity, h is bathymetric elevation, and δ_{clr} is required clearance. Masked, out-of-bounds, or clearance-violating proposals are rejected through the shared constraint path. We report SR, timeout-aware episode time T , $E = \sum_t \|a_t\|_2^2 \Delta t$, near-collision ratio, and task-specific metrics.

BC maps normalized goal, depth, probe, current, and task features ϕ_t to $\hat{a}_t = f_\theta(\phi_t)$ and minimizes the heuristic-action

loss

$$\mathcal{L}_{\text{BC}} = N_{\text{BC}}^{-1} \sum_t \|a_t^{\text{heur}} - \hat{a}_t\|_2^2. \quad (4)$$

Every K steps, an LLM instead proposes a discrete assignment $g_k = \pi_{\text{LLM}}(s_k)$; schema-invalid output invokes a deterministic allocator, and a shared low-level controller executes $a_t = \pi_{\text{low}}(x_t, g_k)$. Appendix Sec. C gives dynamics defaults, pollution semantics, and ROS mapping; Appendix Sec. D defines all metrics and manifests.

Optional pollution fields (OCPNet). Continuous containment applies a local concentration sink. SPC completes discrete sources after an assignment-and-dwell condition; its cleanup rate is the fraction of sources completed. Figure 3 shows the multispecies concentration signal.

Experiments

We compare heuristic, BC, and LLM-assisted controllers across tasks, difficulty tiers, and team sizes, then evaluate task transfer to underwater scenes.

Experimental setup

The main core results use 6DoF with hard mask and clearance checks and a fixed NetCDF-derived current slice; the external backend advances a temporal export in four packaged scenes. We report SR, timeout-aware T , action-energy proxy E , near-collision ratio, and the task metrics named in the tables. Episode budgets accompany each result table. Additional experiments examine planner decisions, dynamics sensitivity, and pollution observations. All runs store their integer seed, task and dynamics configuration, dataset indices, and software and Git state. Local LLM inference assigns one 48 GB accelerator to each process. Appendix Sec. D specifies the supplementary diagnostic budgets and recorded configuration fields.

Here T uses success time on successful episodes and timeout on failures; E accumulates squared clipped vehicle commands over time. Coll(%) counts steps in which the closest agent





Figure 3: OCPNet multispecies diffusion and reaction fields that can generate concentration observations for localization, containment, and cleanup tasks.

Task	N	Heur T2	BC T2	Heur T3	BC T3
<i>Navigation and Tracking</i>					
G2G	1	94.0	44.0	74.0	34.0
SK	1	100.0	0.0	14.0	0.0
RF	1	84.0	24.0	54.0	0.0
DPT	1	84.0	64.0	50.0	0.0
<i>Coordination</i>					
FT	10	100.0	40.0	20.0	0.0
FH	8	94.0	100.0	4.0	4.0
<i>Sensing and Pollution</i>					
PI	1	100.0	90.0	30.0	14.0
AS	1	90.0	84.0	14.0	4.0
SPC	10	84.0	50.0	10.0	0.0
UPL	5	16.0	90.0	20.0	4.0


Table 6: Per-task success rates at Tier 2 and Tier 3 for the 10-task suite, using 50 episodes per task and tier. Highlighted rows are revisited in the planning and external-backend transfer studies.




(a) SPC: Surface Pollution Cleanup (multi-agent)



(b) UPL: Underwater Pollution Lift (5 UAVs)



(c) AS: Area Scan over terrain



(d) FH: Fish Protection / Herding (8 UAVs)

Figure 4: Surface Pollution Cleanup, Underwater Pollution Lift, Area Scan, and Fish Protection / Herding in the core backend.

Method	Avg. SR(%)	Calls	Lat. (s)	ms per call	Prompt tok.	Output tok.	Total tok.
LLM-Qwen2.5-7B	65.0	6.48	14.21	2222.1	3932	4084	8016
LLM-Qwen2-7B	53.3	4.71	11.24	2440.3	2877	2996	5873
LLM-Mistral-7B	53.3	6.05	31.38	4942.5	3762	4128	7890
LLM-ChatGLM3-6B	53.3	5.38	7.91	1393.6	3414	411	3825
Heuristic	55.0	0.00	0.00	0.0	0	0	0
LLM-Llama3-8B	45.0	6.71	13.94	2117.5	3072	3218	6292
LLM-Llama2-7B	33.3	6.59	49.42	7401.8	4219	4851	9071
BC	21.7	0.00	0.00	0.0	0	0	0

Table 7: Core-benchmark planning-suite cost metrics at Tier 2 with doubled currents: uncached local-inference calls, latency, and token usage, reported alongside average success over 20 episodes per task.

pair is within the configured near-collision radius. We also report coverage, inspection, source completion, station error, Cross-Track Error (CTE), and reached waypoints. Each of Tiers 1–3 jointly specifies disturbance, tolerance, spatial extent, and coordination settings.




Figure 5: Success rates across difficulty tiers for a five-task subset. Left: success rate at Tiers 1–3. Right: success-rate drops from Tier 1→2 and Tier 2→3.

Core benchmark backend

Across ten tasks, heuristic success decreases from 84.6% at Tier 2 to 29.0% at Tier 3, while BC decreases from 58.6% to 6.0% (Table 4). Table 6 reports the task-level changes, including gaps in Station Keeping, Route Following, Surface Pollution Cleanup, and Pipeline Inspection. Figure 4 shows four tasks from the suite.

Figure 5 compares success across all three tiers for a five-task subset.

Heuristic control has higher macro success at both tiers. BC performs well on *Underwater Pollution Lift (UPL)* and *Fish Protection / Herding (FH)* at Tier 2, but declines on tasks requiring persistent disturbance rejection or coordinated geometry. In Surface Pollution Cleanup, larger teams improve success and source completion while increasing collision exposure (Figure 6). The action-energy proxy varies non-monotonically with team size.

Planning-sensitive tasks. We compare deterministic heuristic control, BC, and local LLM planners on Surface Pollution Cleanup (SPC), Area Scan (AS), and Pipeline Inspection (PI). Each LLM proposes assignments or waypoints at a fixed stride; a deterministic controller executes them with shared safety checks.

In the 40-episode-per-task study, heuristic control has the highest average success (Table 5). ChatGLM3-6B matches its cleanup rate and achieves the highest scan coverage. In the separate 20-episode-per-task cost study, Qwen2.5-7B has the highest average success, and ChatGLM3-6B has the lowest LLM inference latency (Table 7).

In the three-seed planning diagnostic, valid plans can revisit covered cells or miss leaks, while invalid responses invoke the deterministic task allocator. We report executed command effort through E and inference cost through latency and token counts. Appendix Sec. E reports task-level and model-level validity, fallback, completion, token, and latency values.

External underwater-scene backend

Route-following performance varies across four underwater scenes (Table 8), while Figure 7 shows pollution-response and multi-agent demonstrations. The scene backend introduces scene geometry and applies currents through a one-way disturbance interface. A separate diagnostic found lower localization error with OCPNet-3D under both tested current settings, with




Figure 6: Surface-cleanup scaling at Tier 2. Left: success rate and cleanup rate as the number of vehicles increases. Right: completion time and near-collision ratio, with per-point annotations showing the energy proxy in thousands.



Figure 7: Task demonstrations in packaged underwater scenes using currents exported from the OneOcean dataset.

Scene	Tier 2			Tier 3		
	SR (%)	WP	CTE (m)	SR (%)	WP	CTE (m)
Dam	0.0	1.0	3.06	0.0	0.0	0.87
OpenWater	20.0	3.1	2.50	0.0	0.0	1.91
PierHarbor	56.7	3.5	2.22	0.0	2.0	1.41
SimpleUnderw.	100.0	4.0	3.04	96.7	3.0	2.16
Mean over scenes	44.2	2.92	2.70	24.2	1.25	1.59

Table 8: Route-following across four underwater scenes, with 30 episodes per difficulty in each scene. WP denotes the mean number of reached waypoints; CTE denotes Cross-Track Error.

aggregate success matching or exceeding the analytic probe. Appendix Sec. F provides the values and episode budgets.

Discussion

Heuristic control achieves higher macro success than BC across the ten-task suite, while BC performs well on Underwater Pollution Lift and Fish Protection / Herding at Tier 2. Stronger currents and tighter task requirements expose failures in sustained tracking and coordination. LLM performance varies across tasks and evaluation sweeps, with latency and token usage adding to the cost of planning.

Task-specific progress, command effort, collision exposure, and completion time complement binary success. The planner diagnostics connect task outcomes to plan validity and fallback behavior. Valid assignments can yield poor trajectories or repeated coverage, while invalid calls invoke the deterministic allocator. These observations emphasize the role of closed-loop execution when evaluating high-level planning.

The external experiments show scene-dependent route-following performance under the shared task definitions. Core episodes use fixed current slices, while scene episodes use temporal forcing and scene-specific geometry. Recording these settings supports comparisons across environments and helps identify which conditions challenge a controller.

Conclusion

OneOcean combines a shared ocean dataset, ten marine-robotics tasks, and a standardized evaluation protocol. Its headless and underwater-scene backends support reproducible experiments under data-grounded disturbances. The experiments compare task performance, robustness to currents, planning cost, and transfer across underwater scenes. Future work will improve tide coupling and three-dimensional currents and broaden the tasks and robot morphologies.

References

- Álvarez-Tuñón, O.; Amer, A.; Ugurlu, H. I.; et al. 2023. MIMIR-UW: A Multipurpose Synthetic Dataset for Underwater Navigation and Inspection. Zenodo.
- Chu, S.; Huang, Z.; Li, Y.; Lin, M.; Carlucho, I.; Petillot, Y. R.; and Yang, C. 2025. MarineGym: A High-Performance Reinforcement Learning Platform for Underwater Robotics.
- Copernicus Marine Environment Monitoring Service (CMEMS). 2025. Global Ocean Physics Analysis and Forecast. <https://data.marine.copernicus.eu/>. Accessed: 2025-09-07.
- Copernicus Marine Service. 2024. Copernicus Marine Service: Ocean Monitoring and Forecasting.
- Dagestad, K.-F.; Röhrs, J.; et al. 2018. OpenDrift v1.0: a generic framework for trajectory modelling. *Geoscientific Model Development*, 11(4): 1405–1420.
- Dal Fabbro, N.; Mesbahi, M.; Mendes, R.; Borges de Sousa, J.; and Pappas, G. J. 2026. Long-Term Mapping of the Douro River Plume with Multi-Agent Reinforcement Learning. In *Proc. IEEE International Conference on Robotics and Automation (ICRA)*.
- Delandmeter, P.; and van Sebille, E. 2019. The Parcels v2.0 Lagrangian Framework: New Field Interpolation Schemes. *Geoscientific Model Development*, 12(8): 3571–3584.
- Embley-Riches, J.; Liu, J.; Julier, S.; and Kanoulas, D. 2026. Unreal Robotics Lab: A High-Fidelity Robotics Simulator with Advanced Physics and Rendering. In *Proc. IEEE International Conference on Robotics and Automation (ICRA)*.
- Fu, L.; Yu, J.; El-Refai, K.; Kou, E.; Xue, H.; Huang, H.; Xiao, W.; Wang, G.; Niu, D.; Li, F.-F.; Shi, G.; Wu, J.; Sastry, S.; Zhu, Y.; Goldberg, K.; and Fan, L. 2026. CaP-X: A Framework for Benchmarking and Improving Coding Agents for Robot Manipulation. In *Proc. International Conference on Machine Learning (ICML)*.
- GEBCO Compilation Group. 2025. The GEBCO 2025 Grid – A Global 15 Arc-Second Bathymetry Dataset.
- HYCOM Consortium for Data-Assimilative Ocean Modeling. 2024. HYCOM: Hybrid Coordinate Ocean Model.
- Jangir, Y.; Zhang, Y.; Lo, P.-C.; Yamazaki, K.; Zhang, C.; Tu, K.-H.; Ke, T.-W.; Ke, L.; Bisk, Y.; and Fragkiadaki, K. 2026. RobotArena ∞ : Scalable Robot Benchmarking via Real-to-Sim Translation. In *Proc. International Conference on Learning Representations (ICLR)*.
- Koenig, N.; and Howard, A. 2004. Design and Use Paradigms for Gazebo, An Open-Source Multi-Robot Simulator. In *Proc. IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*.
- Lesser, G. R.; Roelvink, J. A.; van Kester, J. A. T. M.; and Stelling, G. S. 2004. Development and validation of a three-dimensional morphological model. *Coastal Engineering*, 51(8-9): 883–915.
- Lončar, I.; Obradović, J.; Kraševac, N.; Mandić, L.; Kvasić, I.; Ferreira, F.; Slošić, V.; Nađ, Đ.; and Mišković, N. 2022. MARUS - A Marine Robotics Simulator. In *OCEANS 2022, Hampton Roads*.
- Madec, G.; and the NEMO Team. 2019. NEMO Ocean Engine. *Note du Pôle de modélisation de l'Institut Pierre-Simon Laplace (IPSL)*. No. 27.
- Makoviychuk, V.; Wawrzyniak, L.; Guo, Y.; Lu, M.; Storey, K.; Macklin, M.; Hoeller, D.; Rudin, N.; Allshire, A.; Handa, A.; and State, G. 2021. Isaac Gym: High Performance GPU-Based Physics Simulation For Robot Learning. *CoRR*, abs/2108.10470.
- Manhães, M.; Scherer, S.; Voss, M.; Douat, L.; and Stein, F. 2016. UUV Simulator: A Gazebo-based Package for Underwater Intervention and Multi-Robot Simulation. In *Proc. 2016 IEEE/OES Autonomous Underwater Vehicles (AUV)*.
- Prats, M.; Perez, J.; Fernandez, J. J.; and Sanz, P. J. 2012. An Open Source Tool for Simulation and Supervision of Underwater Intervention Missions. In *Proc. 2012 IEEE/RSJ International Conference on Intelligent Robots and Systems*.
- Pshenitsyn, A.; Panov, A.; and Skrynnik, A. 2026. CAMAR: Continuous Actions Multi-Agent Routing. In *Proc. AAAI Conference on Artificial Intelligence*, volume 40, 29651–29659.
- Romrell, B.; Austin, A.; Meyers, B.; et al. 2025. A Preview of HoloOcean 2.0. arXiv:2510.06160.
- Ruediger, M.; and Banerjee, A. G. 2026. Mapping-Guided Task Discovery and Allocation for Robotic Inspection of Underwater Structures. In *Proc. IEEE International Conference on Robotics and Automation (ICRA)*.
- Shchepetkin, A. F.; and McWilliams, J. C. 2005. The Regional Oceanic Modeling System (ROMS): A Split-Explicit, Free-Surface, Topography-Following-Coordinate Oceanic Model. *Ocean Modelling*, 9(4): 347–404.
- Song, J.; Ma, H.; Bagoren, O.; Sethuraman, A. V.; Zhang, Y.; and Skinner, K. A. 2025. OceanSim: A GPU-Accelerated Underwater Robot Perception Simulation Framework.
- Strobach, E.; Molod, A.; Forget, G.; Campin, J.-M.; Hill, C.; Menemenlis, D.; and Heimbach, P. 2018. Consequences of Different Air-Sea Feedbacks on Ocean Using MITgcm and MERRA-2 Forcing: Implications for Coupled Data Assimilation Systems. *Ocean Modelling*, 132: 91–111.
- Wu, Z.; Modi, A.; Mavrogiannis, A.; Joshi, K.; Chopra, N.; Aloimonos, Y.; Karapetyan, N.; Rekleitis, I.; and Lin, X. 2026. DREAM: Domain-Aware Reasoning for Efficient Autonomous Underwater Monitoring. In *Proc. IEEE International Conference on Robotics and Automation (ICRA)*.
- Xue, Y.; Mao, M.; Ru, X.; et al. 2025. OceanGym: A Benchmark Environment for Underwater Embodied Agents. arXiv:2509.26536.
- Zereik, E.; Bibuli, M.; Mišković, N.; Ridao, P.; and Pascoal, A. 2018. Challenges and Future Trends in Marine Robotics. *Annual Reviews in Control*, 46: 350–368.
- Zhang, M. M.; Choi, W.-S.; Herman, J.; Davis, D.; Vogt, C.; McCarrin, M.; Vijay, Y.; Dutia, D.; Lew, W.; Peters, S.; and Bingham, B. 2022. DAVE Aquatic Virtual Environment: Toward a General Underwater Robotics Simulator. arXiv:2209.02862.

Appendix

This appendix details the environment representation, task objectives, motion and pollution models, and evaluation protocol. It also reports complementary studies of dynamics, disturbance strength, planner decisions, and scene-backed multi-agent behavior.

A Ocean Dataset and Validation

Sources, harmonization, and schema

OneOcean combines bathymetry from the General Bathymetric Chart of the Oceans (GEBCO) (GEBCO Compilation Group 2025) with horizontal currents from the Copernicus Marine Service Global Ocean Physics Analysis and Forecast (GOPAF) product (Copernicus Marine Service 2024; Copernicus Marine Environment Monitoring Service (CMEMS) 2025). The resulting Network Common Data Form (NetCDF) product is harmonized by cropping the requested region, retaining returned time and depth coordinates, normalizing coordinate names and units, interpolating variables onto one latitude–longitude grid, preserving invalid terrain and NoData cells, and recording the transformations.

Masked elevation cells use zero as a storage fill value; elevation is interpreted together with the validity mask. In the core simulation frame, $(x, y, z) = (\text{east}, \text{depth-down}, \text{north})$, and the dataset vector (u, v) maps to $(u, 0, v)$. Longitude and latitude are converted to local metric x and z about the recorded lower-left geographic origin.

Depth and tide treatment

The realized `tiny` request contains one level at 0.494 m. The scene and public requests cover 0–200 m and retain 26 source levels from 0.494 to 186.126 m. The pipeline retains the returned depth coordinates. Consequently, `tiny` is a time-varying 2D or 2.5D current product, whereas the other two products provide source-derived horizontal current at multiple discrete depths.

Only the current scene release includes tide components. Hourly tide samples are aligned to the base product with nearest-time selection and attenuated as $\exp(-d/z_0)$ with $z_0 = 50$ m before forming the total current pair. The exported temporal resolution follows the base-product clock; coarse time sampling removes higher-frequency variation.

Dataset validation

The release validator checks required fields, aliases, dimensions, monotonic coordinates, units, finite sampled values, mask semantics, provenance, and optional round-trip fidelity between NetCDF and a NumPy Archive (NPZ). Table A2 summarizes the July 21, 2026 validation of the released products.

The warnings identify missing depth or elevation units, total-current units, or mask flag meanings depending on the variant. Regeneration with the updated pipeline writes these attributes. The tested `tiny` slice preserved coordinates and currents at absolute tolerance 10^{-6} .

Portable provenance

Every regenerated NetCDF records source product identifiers, requested and actual region, time, and depth, tide inclusion and depth profile, interpolation and missing-value handling, and generation software. Companion variant manifests use paths relative to the variant directory. Each run records the selected time and depth indices.

Geographic sampling within an episode

Let (ϕ_0, λ_0) be an episode’s geographic origin and $p = (x, y, z)$ its local position in meters. The core sampler uses

$$\phi(p) = \phi_0 + z/L, \quad (\text{A1})$$

$$\lambda(p) = \lambda_0 + x/[L \max(0.1, \cos \phi_0)], \quad (\text{A2})$$

where $L = 111132$ m per degree and the cosine evaluates latitude in radians. At reset, the origin is sampled within the available geographic extent to place the simulation tile on the environmental grid. The default tile spans 600×600 m horizontally, with vehicle depth bounded between 2 and 18 m. Environmental coordinates retain the source grid spacing.

For grid coordinates $\{\phi_j\}$ and $\{\lambda_k\}$, the sampled indices are

$$j^* = \arg \min_j |\phi_j - \phi(p)|, \quad k^* = \arg \min_k |\lambda_k - \lambda(p)|. \quad (\text{A3})$$

The current at p is $g(u_{j^*k^*}, 0, v_{j^*k^*})$, where g is the current gain. Bathymetry and validity use the same indices. Missing current samples are set to zero, while terrain feasibility is determined separately by the validity mask and seabed clearance. Core episodes query a fixed time and depth slice.

Temporal forcing in the scene backend

The scene backend uses an exported current sequence at a selected geographic location and the nearest available source depth. Let T_c be its number of temporal samples, s the episode seed, and α the configured sequence-advance rate. At simulation time t , the sampled index is

$$i(t) = [(s \bmod T_c) + \lfloor \alpha t \rfloor] \bmod T_c. \quad (\text{A4})$$

Setting $\alpha = 0$ holds the initial sample; a positive value advances the sequence and wraps at its end. For a daily export, α gives source days per simulated second. The sampled vector, multiplied by the current gain, drives the scene disturbance. Spatial slices and temporal sequences provide complementary conditions for evaluating disturbance rejection.

B Task Definitions

Every task declares its objective, observations, feasibility constraints, success condition, and primary metric. Shared observations include pose and relative target cues, local current, and task-specific probes; shared feasibility checks include domain bounds, terrain validity, seafloor clearance, and configured team size.

Go-to-Goal Under Currents (G2G). Reach a goal while rejecting data-grounded drift; success requires entering the goal tolerance, and final goal error is primary.




Figure A1: Conceptual overview of ocean conditions and ecological signals relevant to OneOcean, including terrain, currents, tides, temperature, salinity, weather, biodiversity, and microplastics.

Canonical field	Dimensions	Unit or Encoding	Contract
Latitude	H	degrees north	Strictly monotonic geographic coordinate.
Longitude	W	degrees east	Strictly monotonic geographic coordinate.
Time	T	Climate and Forecast time unit	Monotonic timestamps retained from the source product.
Depth	D	m, positive down	Actual depth coordinates returned by the source.
Eastward and northward current	$T \times D \times H \times W$	m s^{-1}	Base pair (u, v) or tide-combined pair $(u_{\text{total}}, v_{\text{total}})$.
Elevation	$H \times W$	m relative to sea surface	Bathymetric elevation aligned to the current grid.
Validity mask	$H \times W$	binary	Invalid terrain, crop-exterior, or NoData cells.

Table A1: Canonical environment-product schema. The validator accepts documented aliases but maps every accepted product to these semantics.

Variant	Current pair	Errors	Warn.	Round-trip
tiny	uo, vo	0	3	Exact
scene	utotal, vtotal	0	4	Not run
public	uo, vo	0	2	Not run

Table A2: Environment-product validation. All variants are PASS_WITH_WARNINGS; the metadata warnings are listed below. Round-trip denotes the fixed tiny ($t = 0, d = 0$) NetCDF-to-NPZ check.

Station Keeping (SK). Hold position under disturbance; station RMS error measures sustained deviation.

Route Following (RF). Visit an ordered waypoint route; cross-track error and reached waypoints expose partial progress.

Depth-Profile Tracking (DPT). Follow a target depth profile during motion; depth-tracking error is primary.

Area Scan (AS). Cover a bounded region over 3D terrain; grid coverage fraction measures progress.

Pipeline Inspection (PI). Follow a pipeline corridor and encounter leak cues; the inspection score is the detected-leak fraction, with route progress recorded separately.

Formation Transit (FT). Move a team of 2–10 vehicles while maintaining relative geometry; formation error and collision ratio capture coordination quality.

Surface Pollution Cleanup (SPC). Assign 2–10 vehicles to discrete surface sources and satisfy per-source dwell requirements; cleanup rate is completed-source fraction.

Underwater Pollution Lift (UPL). Coordinate five vehicles around a payload and lift it to the surface; lift completion defines success.

Fish Protection / Herding (FH). Use eight vehicles to contain and escort a fish school under currents; centroid progress toward the escort goal defines success.

Exact progress metrics

Several task metrics are normalized ratios with distinct denominators. Let $\mathcal{V}$ be the visited scan cells in an $H \times W$ grid, N_{det} the detected leaks among N_{leak} placed leaks, and N_{done} the sources whose dwell requirement is complete among N_{src}

sources. The released evaluator computes

$$\text{Cov}_{\text{AS}} = \frac{|\mathcal{V}|}{HW}, \quad (\text{A5})$$

$$\text{Score}_{\text{PI}} = \frac{N_{\text{det}}}{N_{\text{leak}}}, \quad (\text{A6})$$

$$\text{Rate}_{\text{SPC}} = \frac{N_{\text{done}}}{N_{\text{src}}}, \quad (\text{A7})$$

$$e_{\text{FT}} = \frac{1}{N} \sum_{i=1}^N \|p_i - (p_g + r_i)\|_2, \quad (\text{A8})$$

$$\text{Prog}_{\text{FH}}(t) = \text{clip}\left(1 - \frac{d_t}{d_0}, 0, 1\right). \quad (\text{A9})$$

Here $p_g + r_i$ is agent i 's assigned formation goal, while d_t and d_0 are the current and initial fish-centroid distances to the escort goal. Each task reports its own progress metric.

Tiers 1, 2, and 3 preserve the task definitions. Each tier specifies a joint setting of current gain, tolerances, route or coverage extent, sensing difficulty, and coordination requirements. Navigation and tracking tasks primarily stress disturbance rejection, sensing and inspection tasks stress spatial coverage and event encounters, and coordination and pollution tasks add team geometry, assignment, or dwell-time coupling.

State-dependent completion

Navigation and station keeping. Let $d_n = \min_i \|p_i^n - p_g\|_2$. G2G completes when d_n enters the goal radius. SK requires consecutive qualifying steps: its hold counter follows

$$h_{n+1} = \begin{cases} h_n + 1, & d_n \leq r, \\ 0, & d_n > r. \end{cases} \quad (\text{A10})$$

Success occurs when h_{n+1} reaches the prescribed hold length. Consequently, an excursion outside the station radius restarts the hold interval. RF and DPT maintain an ordered waypoint index, advance it when a vehicle enters the active waypoint radius, and evaluate completion from the updated index and that step's waypoint distance. DPT additionally requires the team-mean absolute depth error relative to the evaluated waypoint to satisfy its depth tolerance. FT applies its tolerance to every assigned formation goal: $\max_i \|p_i - (p_g + r_i)\|_2 \leq r$.

Area coverage and inspection. AS stores a persistent binary grid of visited cells. With sensor radius r_s and cell width ℓ , a vehicle marks a square neighborhood extending $\lceil r_s/\ell \rceil$ cells in each horizontal direction, clipped at the grid boundary. Cells are counted once. A lawnmower route supplies intermediate goals, while completion depends on the coverage fraction in Equation A5. PI maintains a separate detection flag for each leak. A flag becomes permanent when any vehicle enters the leak's detection radius; the first encounter time is recorded. Detection of all leaks completes the task, and the waypoint index records traversal progress separately.

Accumulated service in surface cleanup. For source j , let a_i^n denote agent i 's assigned source and q the required number

of agents. Its accumulated service time is

$$b_j^n = \sum_i \mathbf{1}[a_i^n = j] \mathbf{1}[\|p_i^n - s_j\|_2 \leq r], \quad (\text{A11})$$

$$\tau_j^{n+1} = \tau_j^n + \Delta t \mathbf{1}[b_j^n \geq q]. \quad (\text{A12})$$

A source completes when τ_j reaches its dwell threshold. Service already accumulated is retained when vehicles leave; this differs from the consecutive hold in SK. Only vehicles assigned to that source contribute to its service count. The task completes once every source is finished. The hard preset couples assignment with rendezvous by requiring two agents to service a source simultaneously.

Attachment and collective lifting. UPL has four phases: approach, lift-off, fifth-agent joining, and ascent. Each of the five agents acquires a persistent attachment flag upon entering the payload attachment radius. Attachment of the first four agents starts the lift-off hold; completing that hold enables joining, and attachment of the fifth agent enables ascent. During the lifting phases the payload follows the attached team's mean position with a phase-dependent depth update. Completion requires the ascent phase, all five attachments, and payload depth at or above the designated surface threshold.

Fish-school escort. FH evolves a school under current drift and a herding response to nearby vehicles. Its progress uses the horizontal distance between the school centroid and escort goal. The eight-vehicle team succeeds when either the normalized progress in Equation A9 reaches its target or the centroid enters the goal radius.

C Implementation Defaults, Pollution, and ROS 2

Dynamics parameters

The main paper defines current coupling, action clipping, diagonal vehicle dynamics, and terrain feasibility. Table A4 lists the Six-Degree-of-Freedom (6DoF) parameters recorded in each run specification.

Each run records whether it uses direct kinematic integration, diagonal 3DoF dynamics, or bounded-attitude 6DoF dynamics.

Discrete motion update and terrain feasibility

The action a_n specifies desired velocity relative to the water in the world frame. Define the speed projection $\Pi_v(a) = a \min(1, v_{\text{max}}/\|a\|_2)$ for nonzero a and $\Pi_v(0) = 0$. In kinematic mode the candidate position is

$$\tilde{p}_{n+1} = p_n + \Delta t[\Pi_v(a_n) + c_n], \quad (\text{A13})$$

where c_n is the sampled world-frame current, including the configured disturbance terms.

For the dynamic modes, let R_n map body vectors to the world frame, $\nu_n = (v_n, \omega_n)$ be body-frame linear and angular velocity, and ω_n^{cmd} the desired angular rate. The latter turns

Task	Tier 2 parameters	Tier 3 parameters	Step budget	
			Tier 2	Tier 3
G2G	Goal radius 4 m	Goal radius 1.3 m	280	300
SK	Radius 2.6 m; hold 110 steps	Radius 1.3 m; hold 240 steps	360	560
RF	12 waypoints; radius 3.2 m	20 waypoints; radius 1.6 m	460	560
DPT	12 waypoints; radius 3.6 m; depth tolerance 1.2 m	18 waypoints; radius 1.8 m; depth tolerance 0.7 m	560	720
FT	Formation radius 36 m; goal tolerance 4 m	Formation radius 70 m; goal tolerance 1 m	600	550
SPC	12 sources; radius 3.8 m; dwell 8 s; one agent	6 sources; radius 2.6 m; dwell 6 s; two agents	720	900
UPL	Attach radius 4 m; hold 3.5 s; surface depth 2.5 m	Attach radius 3 m; hold 4 s; surface depth 2.5 m	720	880
FH	220 fish; progress 0.995 or goal radius 12 m	160 fish; progress 0.99 or goal radius 7 m	650	900
AS	Cell 75 m; sensor radius 35 m; coverage 0.82	Cell 60 m; sensor radius 70 m; coverage 0.95	2000	1700
PI	14 waypoints; 6 leaks; detection radius 5 m	20 waypoints; 8 leaks; detection radius 3.6 m	1500	1500

Table A3: Task-specific parameters for the two principal evaluation tiers. Durations follow from the step budget and simulation timestep; the default Δt is 1 s. UPL uses five vehicles and FH uses eight.

Field	Default	Meaning	Field	Default	Meaning
Δt	1.0 s	Outer step	$v_{\max}$	1.2 m s^{-1}	Relative speed cap
M_ℓ	(12, 12, 12)	Linear inertia proxy	M_a	(6, 6, 6)	Angular inertia proxy
D_ℓ	(8, 8, 8)	Linear damping	D_a	(3, 3, 3)	Angular damping
K_ℓ	(18, 18, 18)	Linear tracking gain	K_a	(10, 10, 10)	Angular tracking gain
K_d	(2, 2, 2)	Angular damping gain	$\omega_{\max}$	(1.2, 1.2, 1.2) rad s^{-1}	Angular-rate cap
k_{yaw}	2.0	Yaw-rate gain	k_{att}	2.0	Roll and pitch stabilization
$\theta_{\max}$	0.6 rad	Roll and pitch limit			

Table A4: Default 6DoF parameters recorded in each run specification.

yaw toward the horizontal action direction and stabilizes roll and pitch toward zero. The force and moment commands are

$$f_n = K_\ell [R_n^\top \Pi_v(a_n) - v_n], \quad (\text{A14})$$

$$m_n = K_a(\omega_n^{\text{cmd}} - \omega_n) - K_d \omega_n. \quad (\text{A15})$$

With the diagonal parameters in Table A4, the velocity update is

$$\tilde{v}_{n+1} = v_n + \Delta t M_\ell^{-1} (f_n - D_\ell v_n), \quad (\text{A16})$$

$$\tilde{\omega}_{n+1} = \omega_n + \Delta t M_a^{-1} (m_n - D_a \omega_n). \quad (\text{A17})$$

Linear speed is projected with Π_v and angular rates are clipped componentwise before advancing position:

$$\tilde{p}_{n+1} = p_n + \Delta t [R_n v_{n+1} + c_n]. \quad (\text{A18})$$

Attitude advances with the clipped angular rates; yaw is wrapped to its principal interval and roll and pitch are bounded by $\theta_{\max}$. The 3DoF mode holds roll and pitch, and their rates, at zero. Position uses the rotation at the start of the step and the updated body velocity.

The candidate position is first clamped to the configured spatial bounds. At this candidate location, valid water requires an admissible mask cell and depth no greater than seabed depth minus the clearance margin, whose default is 1 m. A terrain-infeasible candidate retains the previous position and increments the constraint-violation count; in dynamic modes it also resets relative linear velocity to zero. Near-collision events are measured from pairwise distances after motion. They contribute to the coordination metric without altering the motion update. Task state and completion are then evaluated on the resulting vehicle positions.

Pollution fields and cleanup semantics

The analytic probe field is a current-advected Gaussian,

$$C(p, t) = m(t) C_{\text{peak}} \exp[-\|p - \mu(t)\|^2 / (2\sigma^2)], \quad (\text{A19})$$

with default $\sigma = 35 \text{ m}$ and $C_{\text{peak}} = 1$. An agent in the 35 m continuous-containment neighborhood applies $m(t + \Delta t) = m(t) \max(0, 1 - 0.12\Delta t)$.

Three-dimensional transport and transformation model

The Ocean Current Prediction Network (OCPNet) pollution engine evolves each configured species s through the implemented balance

$$\frac{\partial C_s}{\partial t} = -\mathbf{u} \cdot \nabla C_s + \nabla \cdot (K_s \nabla C_s) - \lambda_s C_s + \sum_r \nu_{sr} R_r + \mathcal{U}_s + \mathcal{P}_s + S_s - Q_s, \quad (\text{A20})$$

$$R_r = k_{r0} \exp[\alpha_r (T - 298.15)] f_r(\text{pH}) \times \prod_{j \in \mathcal{R}_r} C_j, \quad (\text{A21})$$

$$f_r(\text{pH}) = \max(1 + a_r \Delta \text{pH} + b_r \Delta \text{pH}^2, 0).$$

Here $\Delta \text{pH} = \text{pH} - 7$, C_s is concentration, $\mathbf{u} = (u, v, w)$ is velocity, K_s is diffusivity, λ_s is first-order decay, and ν_{sr} is a signed stoichiometric coefficient. The optional terms $\mathcal{U}_s$ and $\mathcal{P}_s$ represent biological uptake and precipitation, while S_s and Q_s collect externally configured sources and sinks.

The code computes spatially varying diffusivity from temperature T , wave velocity U_w , and salinity S as

$$\begin{aligned} K_T &= K_{s0} [1 + 0.02(T - 273.15)], \\ K_w &= K_T + 0.1 U_w, \\ K_s &= K_w [1 + 0.05(1 - S/35)]. \end{aligned} \quad (\text{A22})$$

These coefficients define the diffusivity used in the benchmark. When decay is enabled, its implemented environmental

dependence is

$$\lambda_s = \lambda_{s0} \exp\left(-\frac{E_{a,s}}{RT}\right) f_s(\text{pH}) f_s(O), \quad (\text{A23})$$

$$f_s(\text{pH}) = \max(1 + a_s \Delta \text{pH} + b_s \Delta \text{pH}^2, 0),$$

$$f_s(O) = \max(1 + c_s O + d_s O^2, 0),$$

where the exponential factor is omitted when $E_{a,s} = 0$, R is the gas constant, and O is dissolved oxygen. Biological uptake, when configured, contributes the Michaelis–Menten sink

$$\mathcal{U}_s = -V_{\max,s} e^{\alpha_{T,s}(T-298.15)} (1 + \alpha_{L,s} L) \frac{C_s}{K_{1/2,s} + C_s} B, \quad (\text{A24})$$

with light L and phytoplankton biomass B . For a configured cation–anion precipitation pair, the saturation and rate terms are

$$\text{SI} = \frac{C_M C_A}{K_{\text{sp}}},$$

$$r_p = k_p e^{\alpha_p(T-298.15)} f_p(\text{pH}) \max(\text{SI} - 1, 0), \quad (\text{A25})$$

which subtract r_p from both dissolved species and add it to the configured precipitate before its vertical settling term is applied.

For transport, the implementation uses a first-order upwind difference. Along x , for example,

$$(\delta_x^{\text{up}} C)_i = \begin{cases} (C_i - C_{i-1})/\Delta x, & u_i \geq 0, \\ (C_{i+1} - C_i)/\Delta x, & u_i < 0, \end{cases} \quad (\text{A26})$$

and the centered flux form of the corresponding diffusion term is

$$(\mathcal{D}_x C)_i = \frac{K_i(C_{i+1} - C_i) - K_{i-1}(C_i - C_{i-1})}{\Delta x^2}. \quad (\text{A27})$$

The same constructions are applied along y and z , and explicit Euler combines every active process:

$$C_s^{n+1} = C_s^n + \Delta t \mathcal{F}_s^n, \quad (\text{A28})$$

$$\mathcal{F}_s = \mathcal{A}_s + \mathcal{D}_s - \lambda_s C_s + \sum_r \nu_{sr} R_r$$

$$+ \mathcal{U}_s + \mathcal{P}_s + S_s - Q_s.$$

The solver rejects velocity fields that violate

$$\max\left(\frac{\|u\|_\infty \Delta t}{\Delta x}, \frac{\|v\|_\infty \Delta t}{\Delta y}, \frac{\|w\|_\infty \Delta t}{\Delta z}\right) \leq 1, \quad (\text{A29})$$

and the diffusion module exposes the conservative helper bound $\Delta t_{\text{diff}} = \frac{1}{2} \min(\Delta x, \Delta y, \Delta z)^2 / \|K_s\|_\infty$. Dirichlet, Neumann, periodic, and open boundaries are configurable. The three-dimensional solver leaves negative concentrations unclipped. Timesteps and process parameters must satisfy the stated stability conditions.

The benchmark wrapper uses a $28 \times 28 \times 10$ grid, a 2 s internal step, base diffusivity 8×10^{-8} , point-source emission rate 0.02, 8 m sink radius, and sink strength 0.15 s^{-1} . The benchmark case enables transport, diffusion, the point source, and the

Interface	ROS type and meaning
/agent_NNN/odom	nav_msgs/Odometry; ENU pose and velocity
/agent_NNN/current	Vector3Stamped; ENU current
/agent_NNN/pollution	Float32; local scalar probe
/agent_NNN/bathymetry	Float32; positive water depth
/agent_NNN/cmd_vel	subscribed ENU velocity command
/clock, /episode/metrics	simulation time and JavaScript Object Notation episode state
/reset, /step, /get_task_state	std_srvs/Trigger services

Table A5: ROS 2 interface. Topic prefixes are /oneocean/agents; agent IDs are zero-padded.

robot sink; the other transformation terms are optional. Dataset horizontal current is resampled and repeated over model depth with zero vertical current. The robot sink multiplies cells in its local grid neighborhood by $\max(0, 1 - 0.15 \Delta t)$ at each internal update.

Dataset-driven surface multispecies proxy

The supplementary visualization uses a separate two-dimensional proxy driven by a selected depth slice of the dataset currents. For any field F , one transport substep is

$$\mathcal{T}_n(F) = F + \Delta t [-u_n F_x - v_n F_y + \kappa(F_{xx} + F_{yy})]. \quad (\text{A30})$$

Let A , B , and C denote normalized microplastic, crude-oil, and aggregate fields, respectively. Transport first gives $A^* = \mathcal{T}_n(A^n)$, $B^* = \mathcal{T}_n(B^n)$, and $C^* = \mathcal{T}_n(C^n)$. With $[x]_+ = \max(x, 0)$, aggregation and weathering are then

$$r_a = k_a [A^*]_+ [B^*]_+, \quad r_w = k_w [B^*]_+, \quad (\text{A31})$$

and the implemented update before masking and normalization is

$$\tilde{A} = [A^* - \Delta t r_a]_+ + 0.010 S_A,$$

$$\tilde{B} = [B^* - \Delta t(r_a + r_w)]_+ + 0.008 S_B,$$

$$\tilde{C} = [C^* + \Delta t r_a]_+. \quad (\text{A32})$$

Gaussian source maps S_A and S_B continuously replenish the two parent fields. Land-mask cells are then set to zero, and every nonzero species field is independently divided by its spatial maximum. Figure A2 therefore shows relative concentration patterns and reaction overlap. The surface proxy uses centered numerical gradients; the three-dimensional solver in Equations A20–A29 uses upwind advection.

The SPC presets use 12 sources, 8 s dwell, and one required agent at medium difficulty; 6 sources, 6 s dwell, and two required agents at hard difficulty.

Optional Robot Operating System 2 (ROS 2) Jazzy bridge

The bridge maps core $(x, y, z) = (\text{east}, \text{down}, \text{north})$ to the ROS East-North-Up (ENU) frame $(x, z, -y)$. External commands traverse the same clipping, dynamics, current, constraints, recorder, and task-update path as built-in controllers; internal-control mode can publish observations without accepting external motion commands.

The package was built and launched in a ROS 2 Jazzy container; topic discovery, frame conversion, one externally commanded step, state query, and reset completed successfully. ROS is an optional dependency.




Figure A2: Six sampled times from the dataset-driven surface multispecies proxy. The overlays show normalized microplastic, crude-oil, and aggregate fields together with aggregation and oil-weathering rates. The land mask blocks the proxy fields. Each species is normalized independently by its spatial maximum after every substep.

D Metrics and Reproducibility Protocol

Metric definitions

The main paper defines Success Rate (SR), timeout-aware episode time T , and the action-energy proxy E . The additional team-level aggregate used in the result tables is the near-collision ratio

$$\text{Coll} = \frac{1}{T_{\text{ep}}} \sum_{t=1}^{T_{\text{ep}}} \mathbf{1} \left[\min_{i < j} \|p_t^{(i)} - p_t^{(j)}\|_2 \leq r_{\text{coll}} \right]. \quad (\text{A33})$$

Task-specific metrics follow the definitions in Appendix Sec. B.

Episode and task aggregation

For task k with M_k episodes, let $s_{ke} \in \{0, 1\}$ indicate success and H_{ke} the timeout in seconds. The reported success rate and timeout-aware duration are

$$\text{SR}_k = \frac{1}{M_k} \sum_e s_{ke}, \quad (\text{A34})$$

$$T_k = \frac{1}{M_k} \sum_e [s_{ke} t_{ke}^{\text{succ}} + (1 - s_{ke}) H_{ke}], \quad (\text{A35})$$

where t_{ke}^{succ} is the success time for a completed episode. The success-only duration, when reported separately, averages t^{succ} over successful episodes and is undefined if none succeeds.

An episode’s action-energy proxy sums over both agents and executed steps:

$$E_{ke} = \Delta t \sum_{n=0}^{L_{ke}-1} \sum_{i=1}^N \|\Pi_v(a_{i,n})\|_2^2. \quad (\text{A36})$$

Here L_{ke} is the executed step count. The proxy uses clipped commands. The near-collision ratio above uses one indicator per step, even if several agent pairs are close in that step; tables expressed in percent multiply this ratio by 100.

Station error is accumulated after each motion update. Writing $\bar{d}_n = N^{-1} \sum_i \|p_i^n - p_g\|_2$, its episode RMS is $[L^{-1} \sum_n \bar{d}_n^2]^{1/2}$. The route-tracking diagnostic analogously averages distance to the active waypoint across agents and time. It measures the waypoint-target error associated with the current route stage. For an overall success rate across K tasks, the macro-average is $K^{-1} \sum_k \text{SR}_k$, which gives every task equal weight. With equal episode counts this also equals the pooled episode success rate. Task-specific metrics retain their original units and are averaged within each task and configuration.

Controller observations and decisions

The heuristic, BC, and LLM-assisted controllers share the action interpretation and the motion update in Appendix Sec. C. Heuristics track task targets, BC maps observations to velocity, and the LLM selects discrete task decisions.

Goal-directed control. For an assigned target q_i , the proportional controller produces

$$a_i = \Pi_v(k_p(q_i - p_i)). \quad (\text{A37})$$

The navigation gain is $k_p = 0.9$ and the station-keeping gain is 0.8. Intermediate targets come from the task state: the active waypoint, assigned source, formation offset, or stage of a

Diagnostic	Budget	Configuration	Purpose
Planner diagnostic	3 seeds per task and method	72 episodes; stride 200; 96-token cap	Validity, fallback, completion, and failure diagnosis.
Dynamics diagnostic	20 episodes per task and difficulty	10 tasks; medium and hard; three dynamics modes	Sensitivity to vehicle dynamics.
Current-gain sweep	5 seeds per task and method	300 episodes; hard; 10 tasks; $g \in \{0, 1, 2\}$	Data-current disturbance sensitivity.
Synthetic-tide stress	10 seeds per task and method	200 episodes; hard; 10 tasks; 0.45 m s^{-1} amplitude	Periodic-disturbance sensitivity.
External pollution diagnostic	3 episodes per scene and configuration	Four scenes; two currents; two field models	Localization-model comparison.
External scaling	10 episodes per scene and team size	80 episodes; two scenes; $N \in \{2, 4, 8, 10\}$	Scene-backed team-size sensitivity.

Table A6: Budgets for supplementary diagnostic studies. Headline experiment budgets remain attached to their result tables in the main paper.

coordinated task. This controller corrects positional error at each step under the applied current. The pollution-localization heuristic additionally uses the agents’ scalar probes to estimate a horizontal concentration gradient and coordinate exploration around the highest observed concentration.

Behavior Cloning. The observation for each training example is

$$o_i = [(q_i - p_i)^\top, y_i, C(p_i), c_{i,x}, c_{i,z}, e_k^\top]^\top, \quad (\text{A38})$$

where e_k is a one-hot task indicator. It contains seven continuous features followed by the task encoding. The policy is a multilayer perceptron with two ReLU hidden layers and a three-dimensional linear output. Feature and action means and standard deviations are estimated on the training split. With $\hat{o}$ and $\hat{a}$ denoting standardized values, the learning objective is

$$\mathcal{L}_{\text{BC}}(\theta) = \frac{1}{3M} \sum_{m=1}^M \|f_\theta(\hat{o}_m) - \hat{a}_m\|_2^2. \quad (\text{A39})$$

AdamW minimizes this objective. At inference the output is transformed back to action units and projected onto the permitted speed ball. The reference training configuration uses hidden widths (64, 64), 12 epochs, batch size 4096, learning rate 10^{-3} , zero weight decay, and a randomly selected 10% validation split; the trained checkpoint records its configuration and normalization statistics.

Demonstrations align the recorded command with its task target and local observations. The trajectory converter retains steps with a recorded target. For the kinematic recording convention, which stores the post-step position, it reconstructs the corresponding pre-step position as $p_n = p_{n+1} - \Delta t(a_n + c_n)$ before forming the target displacement.

LLM-assisted planning. At a planning event, the task state is summarized by the task identifier, step, team size, agent positions, and the relevant source or waypoint state. Coordinates are rounded to two decimal places. The planner returns discrete source assignments or waypoint choices, which are checked for admissible indices and task state before use. The task controller then tracks the resulting targets using the same bounded velocity control. Between planning events the current decisions remain active. Failed or invalid responses invoke the task’s deterministic fallback. The diagnostic in Appendix Sec. E separates decision validity, fallback frequency, task completion, and inference cost, to characterize discrete planning within the closed-loop episode.

Hard task	Kinematic SR	3DoF SR	6DoF SR
Formation Transit	55%	50%	50%
Surface Cleanup	95%	85%	85%
Pollution Lift	100%	75%	75%
All-task macro	85.5%	81.5%	81.5%

Table A7: Selected dynamics sensitivity at hard difficulty (20 episodes per task and mode).

Supplementary diagnostic budgets

Table A6 lists the supplementary diagnostic budgets and varied factors. Each study records its tier presets, controller settings, dataset indices, and environment provenance. Dataset validation and ROS interface checks are reported separately.

Compute and run manifests

The main paper reports the compute platform. For provenance, each run stores an integer seed, normalized configuration, task and dynamics specification snapshots, environment-product indices, software and Git state, episode metrics, and recorder streams. Aggregation scripts regenerate comma-separated-value and Markdown result tables from the recorded episode metrics.

Extension priorities

Future extensions include measured three-dimensional flow, vehicle coupling identified through Computational Fluid Dynamics (CFD) or system identification, validated pollutant transport, and ROS-in-the-loop experiments. New controllers, vehicles, and backends use the shared environment queries, declare termination conditions and metrics, and record configurations and dataset indices. They preserve the existing constraint checks and reporting formats.

E Additional Core-Backend Diagnostics

Dynamics-mode sensitivity

Table A7 selects the hard-setting tasks with the largest observed dynamics sensitivity from a 20-episode-per-task diagnostic.

Across all ten hard tasks, macro SR is 85.5% for kinematic and 81.5% for both 3DoF and 6DoF. Under this controller, bounded roll and pitch leave terminal success unchanged in the diagnostic, while direct kinematic integration yields higher macro success.

Current and synthetic-tide stress

Heuristic SR decreases from 48.0% without data-current forcing to 8.0% at doubled current, while its energy and collision exposure increase. BC success varies non-monotonically with current gain and reaches its lowest value at $g = 2$. The

Stress	Method	N	SR	E	Coll.
Data current, $g = 0$	Heuristic	50	48.0	3232.5	5.80
Data current, $g = 1$	Heuristic	50	20.0	3371.6	8.24
Data current, $g = 2$	Heuristic	50	8.0	3775.0	9.86
Data current, $g = 0$	BC	50	6.0	1546.8	4.14
Data current, $g = 1$	BC	50	12.0	1807.1	4.96
Data current, $g = 2$	BC	50	4.0	2323.3	2.44
Synthetic tide	Heuristic	100	17.0	3489.2	6.29
Synthetic tide	BC	100	1.0	1557.6	6.49

Table A8: Hard-setting disturbance diagnostics across all ten tasks. SR and Coll. are percentages. Each data-current row contains five episodes per task. The synthetic-tide rows contain ten episodes per task, use $g = 0$, and apply a 0.45 m s^{-1} sinusoid with a 300 s period.

synthetic-tide stress also separates the methods, with 17.0% SR for the heuristic and 1.0% for BC.

Planner validity, fallback, and cost

The instrumented diagnostic contains three medium-difficulty seeds for each of eight methods and three tasks, for 72 episodes total. Tables A9 and A10 aggregate each Large Language Model (LLM) across its nine episodes. Valid-plan ratio is weighted by planner calls; latency and tokens are weighted by uncached calls.

LLM	SR	Calls	Valid	Fallback	Comp.
ChatGLM3-6B	0.556	49	0.367	0.633	0.758
Llama2-7B	0.333	60	0.950	0.050	0.736
Llama3-8B	0.444	61	0.869	0.131	0.786
Mistral-7B	0.556	55	0.800	0.200	0.804
Qwen2-7B	0.556	43	0.814	0.186	0.790
Qwen2.5-7B	0.667	59	0.864	0.136	0.809

Table A9: Per-model planning outcomes across SPC, AS, and PI (nine episodes per LLM).

LLM	Latency ms per call	Prompt tokens per call	Output tokens per call
ChatGLM3-6B	1466.1	628.1	76.7
Llama2-7B	7429.6	633.7	728.7
Llama3-8B	2068.1	454.1	475.9
Mistral-7B	5150.6	616.5	676.7
Qwen2-7B	2370.6	603.3	628.6
Qwen2.5-7B	2184.5	600.6	623.9

Table A10: Uncached per-call planning cost in the same diagnostic.

Model-level and task-level rankings also diverge: Llama2-7B attains 0.950 validity but only 0.333 SR, whereas ChatGLM3-6B attains 0.367 validity and 0.556 SR. At task level, cleanup’s low valid ratio coexists with high completion because invalid or duplicate assignments can fall back to the deterministic task allocator; AS can revisit covered cells despite valid syntax, and PI still requires physical leak encounters between low-frequency plans. We report validity, fallback frequency, and task success separately.

F Additional External-Backend Diagnostics

The external pollution study compares an analytic advected Gaussian probe with OCPNet-3D in Dam, OpenWater, Pier-Harbor, and SimpleUnderwater. Each combination of scene, model, and current setting uses three localization episodes; Table A12 therefore aggregates 12 episodes per row.

Task	SR	Valid	Fallback	Latency ms per call	Completion
SPC	0.833	0.430	0.570	3199.0	0.972
AS	0.278	0.831	0.169	3602.0	0.758
PI	0.444	0.891	0.109	3337.1	0.611

Table A11: Cross-model planner diagnostics by task (six LLMs, three seeds each).

Scale	Probe	N	SR (%)	Error (m)	Δ error (m)
1.0	Analytic	12	41.7	25.93	0.00
1.0	OCPNet-3D	12	50.0	20.81	−5.12
2.5	Analytic	12	50.0	26.00	0.00
2.5	OCPNet-3D	12	50.0	22.64	−3.36

Table A12: Pollution-localization diagnostic over four scenes (12 episodes per row). Δ error is measured against the analytic probe at the same current scale.

OCPNet-3D lowers mean localization error by 5.12 m at nominal current and 3.36 m under stronger current, while matching or improving aggregate success. Containment quantities are reported separately for each field model: the analytic field tracks a remaining-mass state, while OCPNet tracks grid mass. Table A13 reports removed mass for the analytic-field team-size sweep.

Scene-backed cleanup scaling

Agents	Episodes	SR (%)	Removed mass (kg)
2	20	40.0	0.2081
4	20	100.0	0.4302
8	20	100.0	0.4817
10	20	100.0	0.4886

Table A13: Analytic-containment scaling at nominal current, aggregated over OpenWater and PierHarbor with ten episodes per scene and team size.

Increasing the team from two to four agents raises aggregate success from 40% to 100% and approximately doubles removed mass. The smaller gains from four to ten agents indicate saturation under the fixed 15 s containment horizon.

The route and pollution results show variation across underwater scenes. Each external-backend run records the scene, current scale, dataset clock, controller settings, and observation model.